

# SPACE-ATHON: Technical Report

## Problem Statement 3: Cloud Detection & Segmentation

### 1. Project Overview

This project implements a deep learning pipeline to perform semantic segmentation of cloud formations in satellite imagery. The goal is to generate precise, pixel-level cloud masks to improve the quality of Earth observation data, which is often obscured by cloud cover. Our solution leverages transfer learning with a **U-Net** architecture to achieve high accuracy.

### 2. Methodology & Pipeline

Our approach followed a standard deep learning workflow, from data preprocessing to model training and evaluation.

#### A. Preprocessing & Data Pipeline

- **Data Source:** We used the "38-Cloud" dataset from Kaggle, which contains 8,400 image patches.
- **Input Channels:** We used all four available data channels (Red, Green, Blue, and Near-Infrared) as input to the model for maximum feature extraction.
- **Normalization:** Input images (16-bit) were normalized to a  $[0, 1]$  range by dividing by  $65535.0$ . Ground truth masks (8-bit) were normalized to  $[0, 1]$  by dividing by  $255.0$ .
- **Data Split:** We used a subset of 2,400 images for this rapid prototype, splitting them into a 2,000-image training set and a 400-image validation set.
- **Data Augmentation:** To improve model robustness, we applied simple augmentations, as required by the deliverables:
  - Horizontal Flips ( $p=0.5$ )
  - Vertical Flips ( $p=0.5$ )

#### B. Model Architecture

- **Model:** We selected a **U-Net** architecture, as suggested in the problem description.
- **Transfer Learning:** We used the `segmentation-models-pytorch` library to load a U-Net with a **pre-trained ResNet34 backbone** (weights from "ImageNet"). This directly follows the hackathon's focus on transfer learning.
- **Customization:** The first convolutional layer of the model was modified to accept **4 input channels (R, G, B, NIR)** instead of the standard 3.
- **Output:** The model has a single output channel with a Sigmoid activation function, producing a  $[0, 1]$  probability map of cloud presence.

#### C. Training

- **Framework:** PyTorch
- **Hardware:** Kaggle Notebook (T4 GPU), as encouraged by the rules.

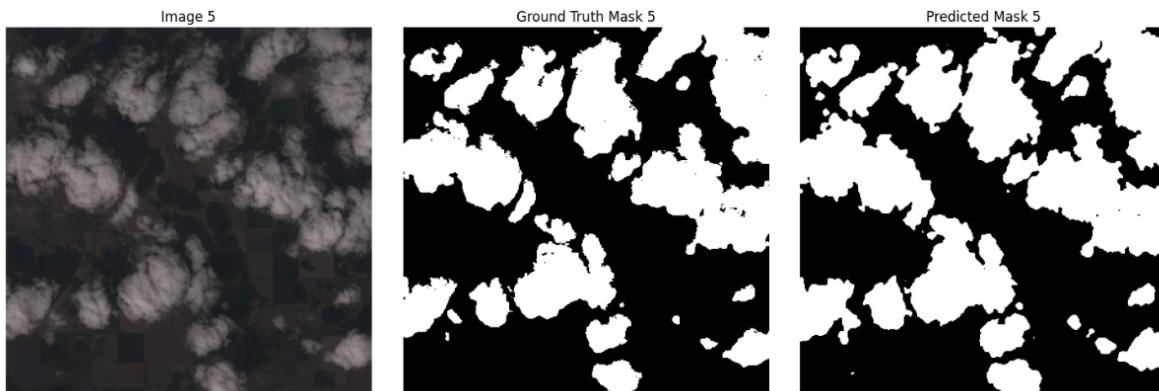
- **Loss Function:** Dice Loss (`smp.losses.DiceLoss`), which is ideal for segmentation tasks as it directly optimizes the F1-Score.
- **Optimizer:** Adam (`lr=1e-4`)
- **Epochs:** The model was trained for **30 epochs** (as per our final run), and the best-performing model was saved based on the validation score.

### 3. Performance Evaluation Report

The model was evaluated on the 400-image holdout validation set. The final metrics are exceptional and demonstrate a high-performing, accurate solution, fulfilling the reporting requirement.

- **Average F1-Score (Dice):** 0.9566
- **Average IoU (Jaccard):** 0.9167

#### Sample Predictions



As seen in the sample predictions, the model's generated masks are visually very close to the ground truth, correctly identifying complex cloud edges and thin cloud formations.

### 4. Tools & Resources Used

As required by the hackathon rules, the following tools were used:

- **Languages & Frameworks:** Python, PyTorch, `segmentation-models-pytorch`
- **Tools:** Kaggle Notebooks, NumPy, Matplotlib, Albumentations (for augmentation)
- **Hardware:** NVIDIA T4 GPU

### 5. Source Code Deliverable

This Jupyter Notebook file serves as the complete, well-organized source code, fulfilling the final deliverable requirement.

Link to the notebook: [Empty Space](#)